

Start coding or [generate](#) with AI.

MILEAGE PREDITON - REGRESSION ANALYSIS

Source:

This dataset was taken from the StatLib library which is maintained at Carnegie Moellon University. The dataset was used in the 1983 American Statistical Association Exposition.

Data Set Information:

This dataset is a slightly version of the dataset provided in the StatLib library. In line with the use by Ross Quinlan 1993 in predicting the attribute "MPG", 8 if the original were removed because they had unknown values for the "MPG" attribute. The original dataset is available in the file "auto-mpg.data-original". "The data concerns city-cycle fuel consumption in miles per gallon, to be predicted in terms of 3 multivalued discrete and 5 continuous attributes."

Import library

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

Import data

```
mp=pd.read_csv('https://github.com/YBI-Foundation/Dataset/raw/main/MPG.csv')
```

```
mp.head()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	origin	name
0	18.0	8	307.0	130.0	3504	12.0	70	usa	chevrolet chevelle malibu
1	15.0	8	350.0	165.0	3693	11.5	70	usa	buick skylark 320
2	18.0	8	318.0	150.0	3436	11.0	70	usa	plymouth satellite
3	16.0	8	304.0	150.0	3433	12.0	70	usa	amc rebel sst
4	17.0	8	302.0	140.0	3449	10.5	70	usa	ford torino

```
mp.nunique()
```

```
mpg      129
cylinders    5
displacement  82
horsepower   93
weight     351
acceleration  95
model_year   13
origin        3
name       305
dtype: int64
```

Data preprocessing

```
mp.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg             398 non-null   float64
1   cylinders       398 non-null   int64
2   displacement    398 non-null   float64
```

```

3  horsepower    392 non-null    float64
4  weight        398 non-null    int64
5  acceleration  398 non-null    float64
6  model_year    398 non-null    int64
7  origin        398 non-null    object
8  name          398 non-null    object
dtypes: float64(4), int64(3), object(2)
memory usage: 28.1+ KB

```

```
mp.describe()
```



	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year
count	398.000000	398.000000	398.000000	392.000000	398.000000	398.000000	398.000000
mean	23.514573	5.454774	193.425879	104.469388	2970.424623	15.568090	76.010050
std	7.815984	1.701004	104.269838	38.491160	846.841774	2.757689	3.697627
min	9.000000	3.000000	68.000000	46.000000	1613.000000	8.000000	70.000000
25%	17.500000	4.000000	104.250000	75.000000	2223.750000	13.825000	73.000000
50%	23.000000	4.000000	148.500000	93.500000	2803.500000	15.500000	76.000000
75%	29.000000	8.000000	262.000000	126.000000	3608.000000	17.175000	79.000000
max	46.600000	8.000000	455.000000	230.000000	5140.000000	24.800000	82.000000

Start coding or [generate](#) with AI.

✓ Removing Missing Values

```
mp=mp.dropna()
```

```
mp.info()
```



```

<class 'pandas.core.frame.DataFrame'>
Index: 392 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   mpg              392 non-null    float64
1   cylinders        392 non-null    int64
2   displacement     392 non-null    float64
3   horsepower       392 non-null    float64
4   weight           392 non-null    int64
5   acceleration     392 non-null    float64
6   model_year       392 non-null    int64
7   origin           392 non-null    object
8   name             392 non-null    object
dtypes: float64(4), int64(3), object(2)
memory usage: 30.6+ KB


```

✓ Data Visualization

```
sns.pairplot(mp,x_vars=['displacement','horsepower','weight','acceleration'],'mpg',y_vars=['mpg'])
```




```
<seaborn.axisgrid.PairGrid at 0x7cd368ec7af0>
```



```
sns.regplot(x='displacement',y='mpg',data=mp)
```

```
<Axes: xlabel='displacement', ylabel='mpg'>
```



✓ Define target variable y and feature x

```
mp.columns
```

```
Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
       'acceleration', 'model_year', 'origin', 'name'],
      dtype='object')
```

```
y=mp['mpg']
```

```
x=mp[['displacement', 'horsepower', 'weight', 'acceleration']]
```

```
y.shape
```

```
(392,)
```

```
x.shape
```

```
(392, 4)
```

```
x
```

```
displacement  horsepower  weight  acceleration
0           307.0         130.0   3504           12.0
1           350.0         165.0   3693           11.5
2           318.0         150.0   3436           11.0
3           304.0         150.0   3433           12.0
4           302.0         140.0   3449           10.5
...          ...          ...     ...           ...
393          140.0          86.0   2790           15.6
394           97.0          52.0   2130           24.6
395          135.0          84.0   2295           11.6
396          120.0          79.0   2625           18.6
397          119.0          82.0   2720           19.4
```

```
392 rows × 4 columns
```

```
y
```

```
0      18.0
1      15.0
2      18.0
3      16.0
4      17.0
```

```

...
393 27.0
394 44.0
395 32.0
396 28.0
397 31.0
Name: mpg, Length: 392, dtype: float64

```

✓ Scaling data

```
from sklearn.preprocessing import StandardScaler
```

```
ss=StandardScaler()
```

```
x=ss.fit_transform(x)
```

```
x
```

```

array([[ 1.07728956,  0.66413273,  0.62054034, -1.285258  ],
       [ 1.48873169,  1.57459447,  0.84333403, -1.46672362],
       [ 1.1825422 ,  1.18439658,  0.54038176, -1.64818924],
       ...,
       [-0.56847897, -0.53247413, -0.80463202, -1.4304305  ],
       [-0.7120053 , -0.66254009, -0.41562716,  1.11008813],
       [-0.72157372, -0.58450051, -0.30364091,  1.40043312]])

```

```
mp=pd.DataFrame(x)
```

```
mp.describe()
```

```

count    0          1          2          3
count  3.920000e+02  3.920000e+02  3.920000e+02  3.920000e+02
mean   -7.250436e-17 -1.812609e-16 -1.812609e-17  4.350262e-16
std     1.001278e+00  1.001278e+00  1.001278e+00  1.001278e+00
min    -1.209563e+00 -1.520975e+00 -1.608575e+00 -2.736983e+00
25%    -8.555316e-01 -7.665929e-01 -8.868535e-01 -6.410551e-01
50%    -4.153842e-01 -2.853488e-01 -2.052109e-01 -1.499869e-02
75%     7.782764e-01  5.600800e-01  7.510927e-01  5.384714e-01
max     2.493416e+00  3.265452e+00  2.549061e+00  3.360262e+00

```

✓ Train Test split data

[+ Code](#)
[+ Text](#)

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.7,random_state=(2529))
```

```
x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
((274, 4), (118, 4), (274,), (118,))
```

✓ Linear Regression model

```
from sklearn.linear_model import LinearRegression
```

```
lr=LinearRegression()
```

```
lr.fit(x_train,y_train)
```

```

LinearRegression()

```

```
lr.intercept_
```

```
↳ 23.485738559737584
```

```
lr.coef_
```

```
↳ array([-1.05767743, -1.68734727, -4.10787617, -0.11495177])
```

✓ Predict test data

```
y_pred=lr.predict(x_test)
```

```
y_pred
```

```
↳ array([18.51865637, 15.09305675, 14.30128789, 23.6753321 , 29.7546115 ,  
23.68796629, 26.61066644, 24.56692437, 15.06260986, 11.94312046,  
24.08050053, 27.96518468, 31.66130278, 31.01309132, 18.32428976,  
19.32795009, 28.08847536, 32.1506879 , 31.15859692, 27.15792144,  
18.82433097, 22.54580176, 26.15598115, 32.36393869, 20.74377679,  
8.78027518, 22.19699435, 18.20614294, 25.00052718, 15.26421552,  
23.13441082, 17.10542257, 9.87180062, 30.00790415, 20.41204655,  
29.11860245, 24.4305187 , 21.72601835, 10.51174626, 13.12426391,  
21.41938406, 19.96113872, 6.19146626, 17.79025345, 22.5493033 ,  
29.34765021, 13.4861847 , 25.88852083, 29.40406946, 22.41841964,  
22.07684766, 16.46575802, 24.06290693, 30.12890046, 10.11318121,  
9.85011438, 28.07543852, 23.41426617, 20.08501128, 30.68234133,  
20.92026393, 26.78370281, 22.9078744 , 14.15936872, 24.6439883 ,  
26.95515832, 15.25709393, 24.11272087, 30.80980589, 14.9770217 ,  
27.67836372, 24.2372919 , 10.92177228, 30.22858779, 30.88687365,  
27.33992044, 31.18447082, 10.8873597 , 27.63510608, 16.49231363,  
25.63229888, 29.49776285, 14.90393439, 32.78670687, 30.37325244,  
30.9262743 , 14.71702373, 27.09633246, 26.69933806, 29.06424799,  
32.45810182, 29.44846898, 31.61239999, 31.57891837, 21.46542321,  
31.76739191, 26.28605476, 28.96419915, 31.09628395, 24.80549594,  
18.76490961, 23.28043777, 23.04466919, 22.14143162, 15.95854367,  
28.62870918, 25.58809869, 11.4040908 , 25.73334842, 30.83500051,  
21.94176255, 15.34532941, 30.37399213, 28.7620624 , 29.3639931 ,  
29.10476703, 20.44662365, 28.11466839])
```

✓ Model Accuracy

```
from sklearn.metrics import mean_absolute_error,mean_absolute_percentage_error,r2_score
```

```
mean_absolute_error(y_test,y_pred)
```

```
↳ 3.3286968643244106
```

```
mean_absolute_percentage_error(y_test,y_pred)
```

```
↳ 0.14713035779536746
```

```
r2_score(y_test,y_pred)
```

```
↳ 0.7031250746717691
```